

# GumGum ML Engineer take-home project

## Overview

GumGum is launching a new promotional website that allows anyone to upload an image and get an instant classification of their image using a pre-trained model capable of identifying 1000 classes of objects. GumGum wants to go live in a week but the front-end (which is being built by a separate team) needs an API to power it. This is exactly where you, our fancy new Machine Learning Engineer, come in to save the day!

### Deadline

This project must be submitted no later than a week after the time these requirements were delivered to your email.

## Classification model

### Model

Download the pre-trained MobileNet classifier [here](#)

### Classification vector indices

Download [here](#)

### Pixel values

Pixel values for input images need to be normalized between -1 and 1 (example: 0,255 (ints) should be scaled to -1,1 (floats) and 0, 127 to -1, 0). If your input image has an alpha channel, you can safely remove/ignore it.

### Resolution

Images loaded into the model need to be resized to 224px by 224px (224, 224, 3) regardless of original dimensions.

### Encoding

We want to support input images in the JPEG, PNG and GIF formats.

## API requirements

Using the language, frameworks, libraries, data stores and tooling of your choice, build a REST API that meets the following requirements:

### 1) Image classifier endpoint

Using TensorFlow , integrate the classification model described in the previous section with your

API. Then, create the endpoint `POST /classify-image` that accepts a single JSON parameter `image_url` which contains a link to the image to be classified. Your API should download the image, re-size and re-map it as required by the classification model and return a response that indicates the most-likely classification, the confidence (in the range of 0-1) of that classification and the `processing_time` (s) required to process and classify the image.

#### Example request

```
curl --request POST --header "Content-Type: application/json" --data
'{"image_url": "https://s3.amazonaws.com/gumgum-interviews/ml-engineer/cat.jpg"}'
http://localhost/classify-image
```

#### Example response

```
{
  "classification": "tabby",
  "confidence": "0.554",
  "processing_time": "4.731"
}
```

Note: this **is** the expected response for the example image above.

## 2) Classification caching

Now that your API can classify external images, improve your endpoint's processing time by introducing server-side caching based on the `image_url`. When implemented, any subsequent requests of the same URL should return the cached classification. To avoid serving stale classifications, be sure to set the TTL on the cache keys to 1 hour.

## 3) Reporting endpoint

Now that we can serve classifications in a performant manner, we can start improving the visibility of the service's usage via a new reporting endpoint `GET /report`. When hit, the endpoint should return a list of the top 10 most requested `image_urls` along with the following metadata for each:

```
Number of times the image was classified (including cached responses)
The minimum, maximum and average total processing times
```

## Bonus

### 4) Dockerize your API

Write a Dockerfile to build a Docker image containing everything needed to run your API service. Once completed, we should be able to start your service with a simple `docker build -t classifier-api .` followed by a `docker run classifier-api`. Additional build/run arguments are allowed but should be documented in your project's `readme.md` file.

## 5) Elasticsearch monitoring

ElasticSearch is great for storing and querying time-series data. GumGum would like to be able to query production data for your new service in a real-time fashion using Kibana, a popular web UI that sits on top of ElasticSearch. First, spin up an AWS Elasticsearch Service cluster comprising of a single node using the AWS Console (no need to automate this). Doing so will also provision a Kibana instance that is pre-configured to connect to your cluster.

Once you have verified you can connect to your ElasticSearch service, update your API to create a new ElasticSearch document for each request to your GET /report endpoint. The document should include all the fields already returned by your API.

Then, create a simple visualization of your real-time data in Kibana and include a screenshot of it in your project's `readme.md`.

## 6) User authentication

Using the user management pattern of your choice, create a new endpoint GET /secret that returns Hi <username> for authenticated users only (otherwise return a 401 error).

# Questions

### Question type

### Contact

Machine learning [mgreenberg@gumgum.com](mailto:mgreenberg@gumgum.com)

API development [corey@gumgum.com](mailto:corey@gumgum.com)

And of course feel free to contact either of us for general questions or comments :)

## How to submit your project

Write a basic `readme.md` for your project that documents how to build and run your API server in a `localhost` fashion. Please include an example `curl` request and response to your POST /classify-image endpoint. Compress all your files into a single `.tar.gz` or `.zip` and send it to the following GumGummers:

[bfuller@gumgum.com](mailto:bfuller@gumgum.com)  
[cambron@gumgum.com](mailto:cambron@gumgum.com)  
[corey@gumgum.com](mailto:corey@gumgum.com)